

**CSci 5451: Parallel Computing–Spring 2011**

# **Homework 4: CUDA Longest Common Subsequence**

Harlan Iverson

April 13, 2011

The solution is not working code. The idea was to create a grid with blocks and have them computed in a diagonal order, from top left to bottom right. Each block depends on the ones directly above and to the left being computed. In order to preserve spatial/temporal locality, a mapping that stores rows of the grid in a diagonal-major order was used.

The longer string should be along the y axis, and the shorter string along the x axis. This would allow a maximum number of concurrent blocks for the most steps, and help spatial locality. The first and last steps would be a single block; and each step from the beginning would have incrementally more until a maximum is reached, and then each step toward the end would have incrementally less.

The code is semi-complete and the algorithm works in serial using a for loop that 'simulates' the device by looping over x values (and uses a 1x1 block configuration, since a thread would be in charge of a single column). The working implementation is in `lcs_serial.cu`. The non-working partial CUDA port is in `lcs_cuda.cu`. Some of the calculation code is in `calc.c`.

The procedure of the calculation is as follows:

- Create a grid of blocks(maxb, maxt, width, height)
  - ensure that  $\text{maxt} * \text{maxb} \leq \text{width}$
  - block width = maxt
  - block height = some constant
  - # y blocks =  $\text{ceil}(h / \text{block height})$
  - # x blocks =  $\text{ceil}(w / \text{block width})$
  - # of diagonals = # x blocks + # y blocks - block width
- for each diagonal (d = 0 to # of diagonals - 1)
  - start block =  $\text{max}(d - \text{\# x blocks} - 1, 0)$

- end block =  $\min(d, \# \text{ x blocks} - b)$
- # blocks = end block - start block
- run CUDA with  $\lll \# \text{ blocks}, \text{maxt} \ggg (\text{grid}, \text{start block}, d, \text{block height})$
- \* block x = (start block + cuda.blockIdx)\*cuda.blockDim [\* block width is covered by blockDim]
- \* block y =  $((d+1) * \text{block height}) - \text{block x}$
- \* block x += 1, block y += 1
- \* for y = block y to block y + block height
  - do LCS cell procedure on (x,y)

Trace back is identical to the serial counter-part.